Map Reduce and Streaming Calculations

Giri Iyengar

Cornell University

gi43@cornell.edu

April 16, 2018

Agenda for the week

- Map-Reduce
- Poisson resampling
- Streaming Calculations
  1. Reservoir Sampling
  2. Storing Items in Sets
  3. Counting in single pass
  4. Frequent Items in a stream
  5. Estimating CDF/PDF in streaming mode
- Background Reading

Overview

Hello World in Map Reduce

Counting words in some text

But I must explain to you how all this mistaken idea of denouncing pleasure and praising pain was born and I will give you a complete account of the system.

- Step 0: Parse text word by word.
- Step 1: Emit one word at a time
- Step 2: Group same words together
- Step 3: Count occurrence of each word

Word Count - Emit

```
But                1
I                1
must                1
explain                1
to                1
you                1
how                1
all                1
this                1
mistaken                1
idea                1
of                1
denouncing                1
pleasure                1
```

Word Count - Emit

```
and                 1
praising                1
pain                 1
was                 1
born                 1
and                 1
I               1
will                 1
give                 1
you                 1
a               1
complete                1
account                 1
of                1
the                 1
system.                 1
```

Word Count - Sort and Group

```
But                    1
I                  1
I                  1
a                  1
account                1
all                    1
and                    1
and                    1
born                    1
complete                1
denouncing                 1
explain                1
give                    1
how                    1
idea                    1
mistaken                1
```

Word Count - Sort and Group

```
must                1
of              1
of              1
pain                1
pleasure                1
praising                1
system.                 1
the              1
this                1
to              1
was              1
will                1
you              1
you              1
```

Word Count - Aggregate

```
But                  1
I                2
a                1
account                  1
all                  1
and                  2
born                   1
complete                     1
denouncing                       1
explain                  1
give                   1
how                  1
idea                   1
mistaken                     1
```

Map Reduce

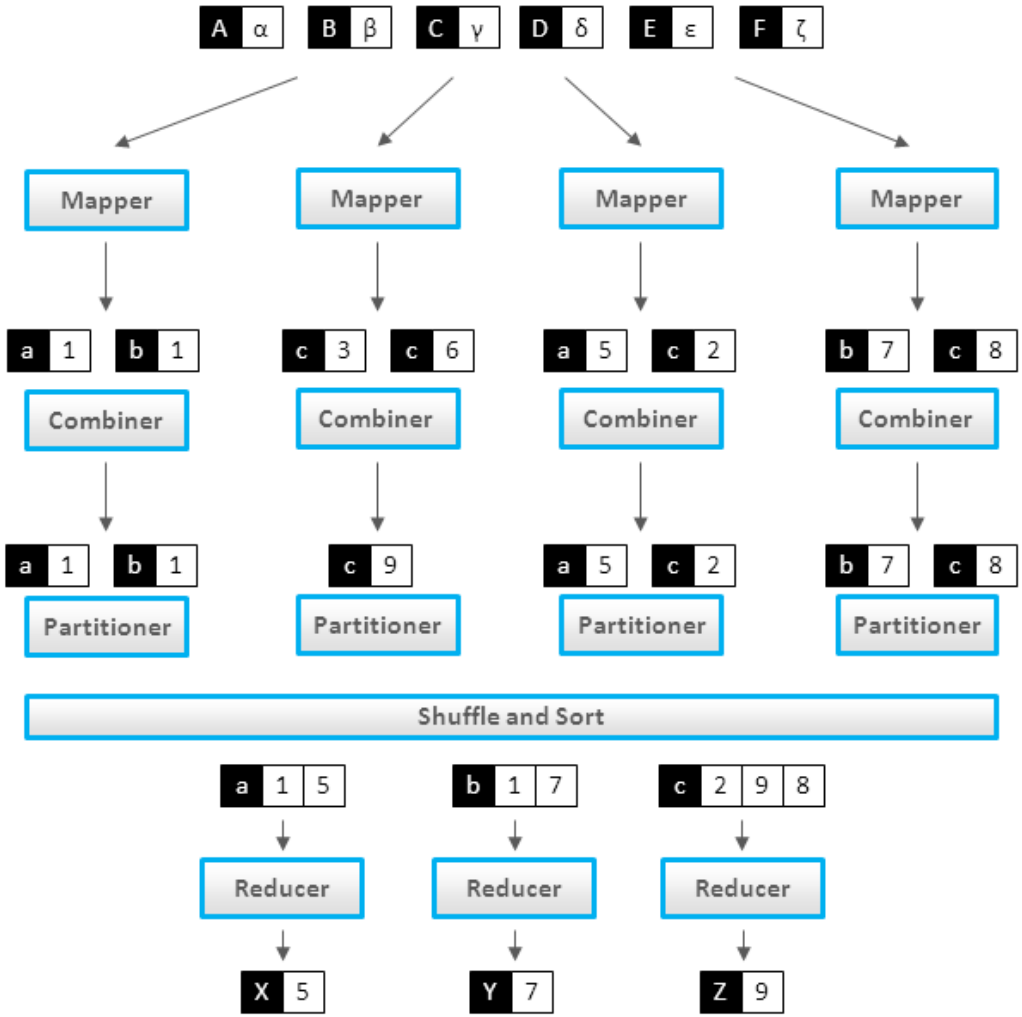

Figure: Source: highlyscalable.com blog

MR Demo

Map Reduce Examples

- Word Count

Map Reduce Examples

- Word Count
- Unique Count

Map Reduce Examples

- Word Count
- Unique Count
- Total Sales, Average Sales by Customer

Map Reduce Examples

- Word Count
- Unique Count
- Total Sales, Average Sales by Customer
- Click-Through-Rate of Advertising Campaigns

Map Reduce Examples

- Word Count
- Unique Count
- Total Sales, Average Sales by Customer
- Click-Through-Rate of Advertising Campaigns
- Little Bag of Bootstraps to Build Models

HDFS: Hadoop Distributed File System

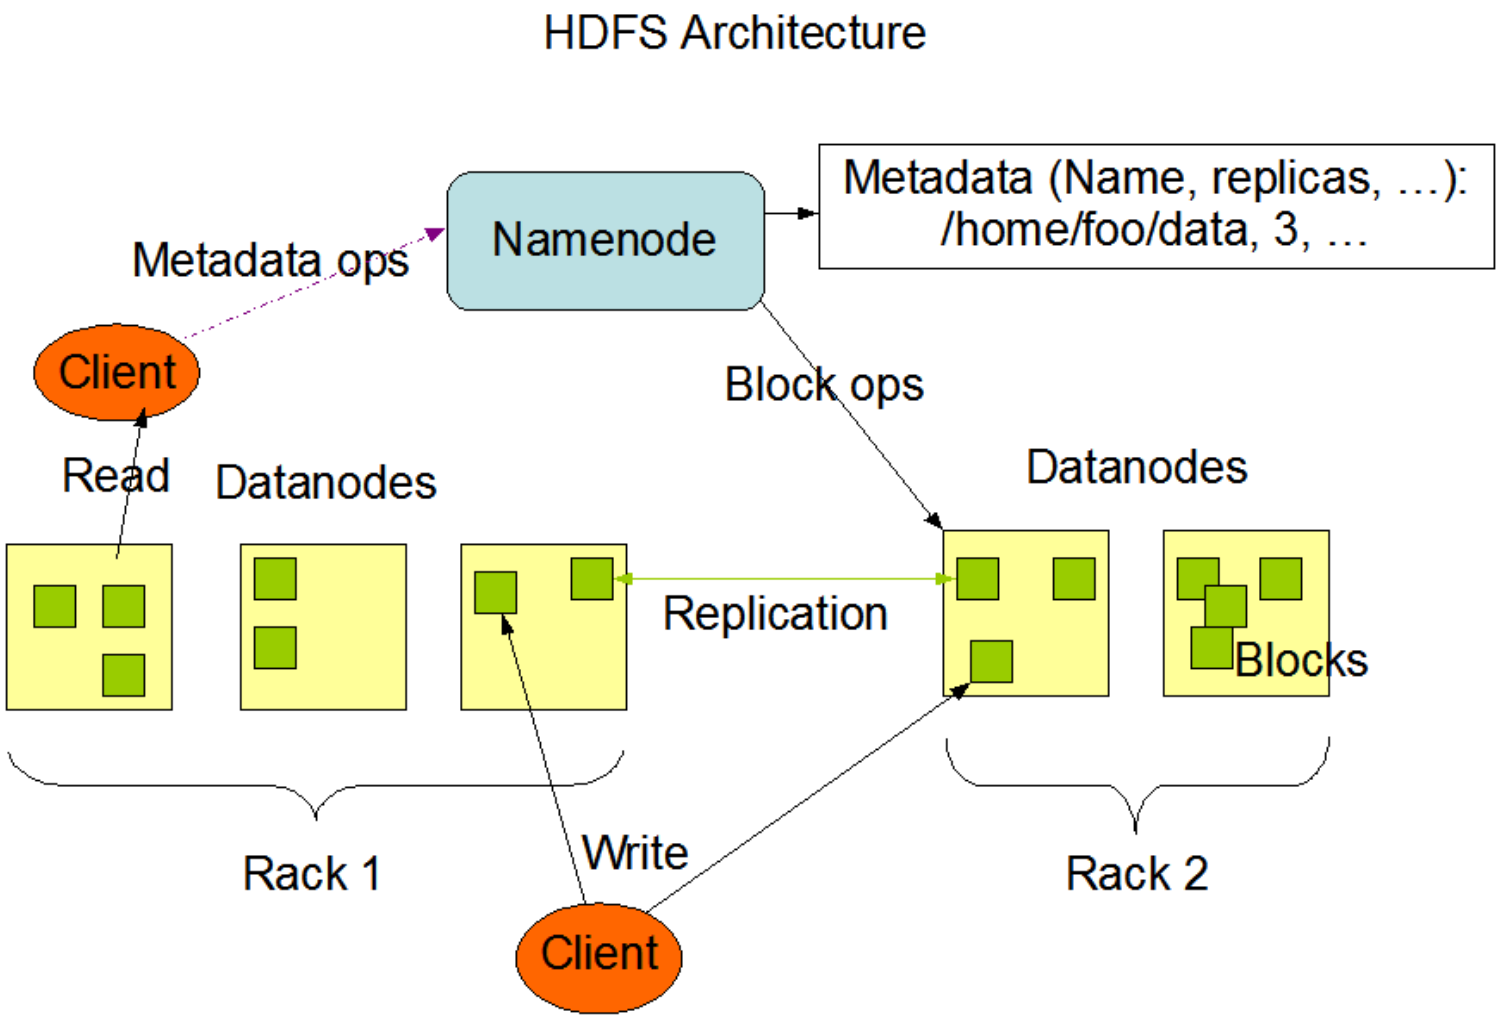

Figure: Source: Apache Foundation

Hadoop Ecosystem

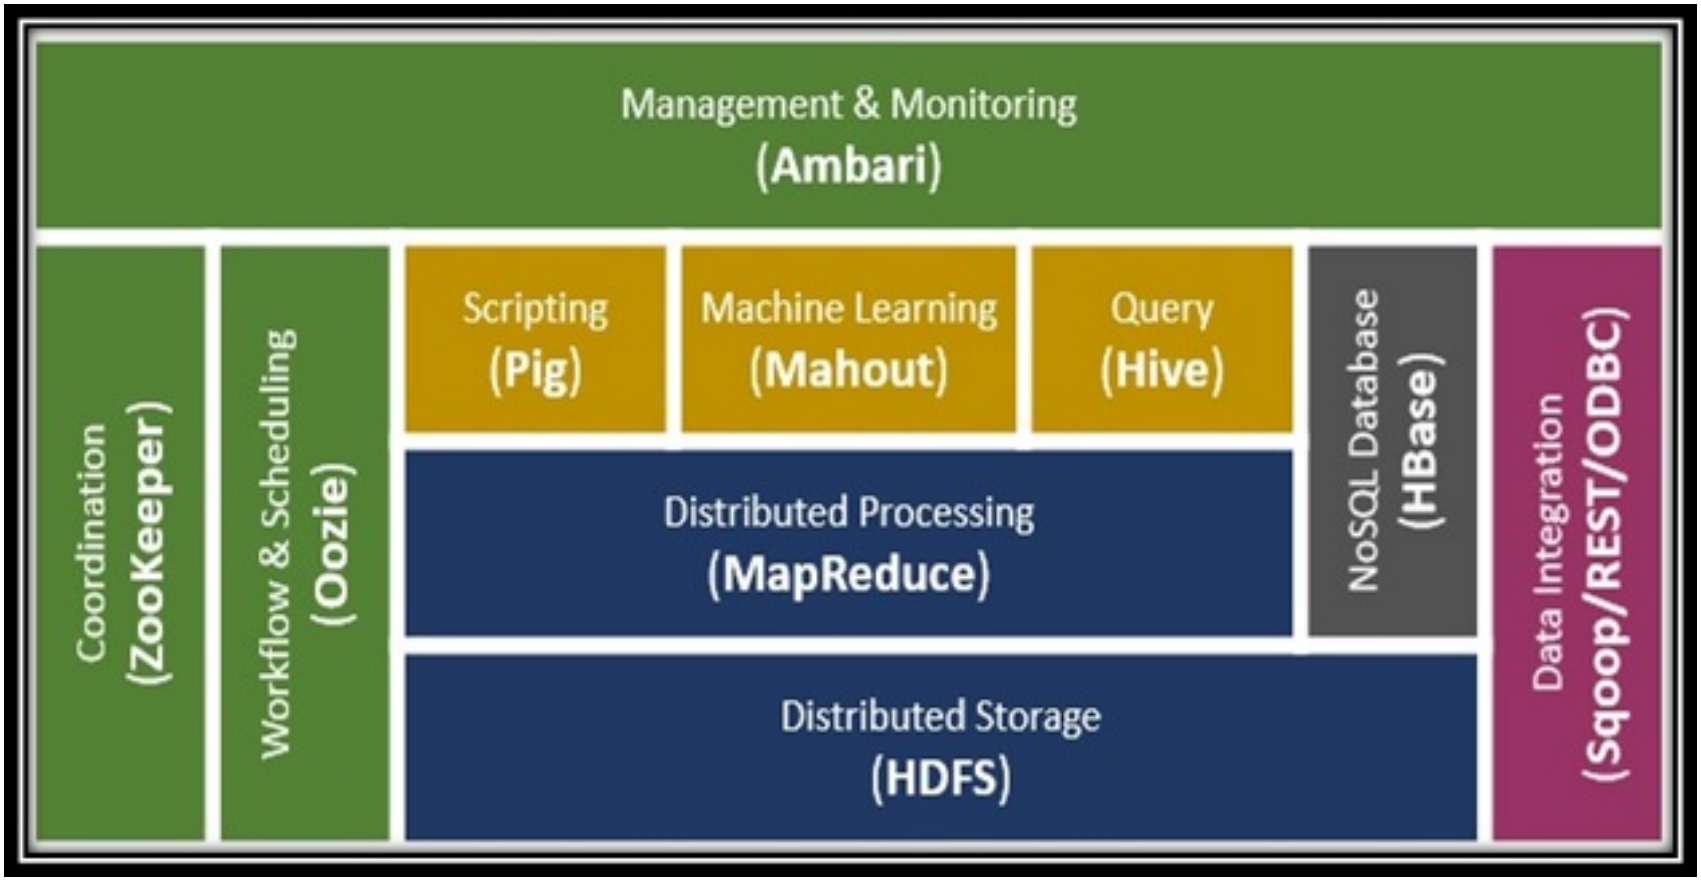

Figure: Source: Apache Foundation

Overview

Poisson Resampling for Efficient Bootstrapping

Bootstrapping Big Data

- Bootstrap is not efficient when dealing with Big Data $\approx 63.2\%$ of data gets resampled per bootstrap
- Need several bootstrap samples depending on what you are trying to estimate (e.g. std err or percentiles)
- Little Bag of Bootstraps is one technique (saw last week)
- Poisson Resampling is another technique

Poisson Resampling

Motivation

Let's start with a tiny sample $\{1.5, 2.5, 3.5, 4.5\}$ and do bootstrap

Sample 1	$\{1.5, 1.5, 3.5, 2.5\}$
Sample 2	$\{2.5, 1.5, 3.5, 2.5\}$
Sample 3	$\{3.5, 4.5, 4.5, 4.5\}$

We can actually describe this in terms of sample counts

Sample 1	$\{2, 1, 1, 0\}$
Sample 2	$\{1, 2, 1, 0\}$
Sample 3	$\{0, 0, 1, 3\}$

These counts follow a $Multinomial(4, \frac{1}{4}, \frac{1}{4}, \frac{1}{4}, \frac{1}{4})$ distribution. Generally, $Multinomial(n, \frac{1}{n}, \ldots, \frac{1}{n})$

Poisson Resampling

Big Data Problem

Not all data resides in the same place. Data is distributed. Also, in streaming cases, we may not even know n in advance

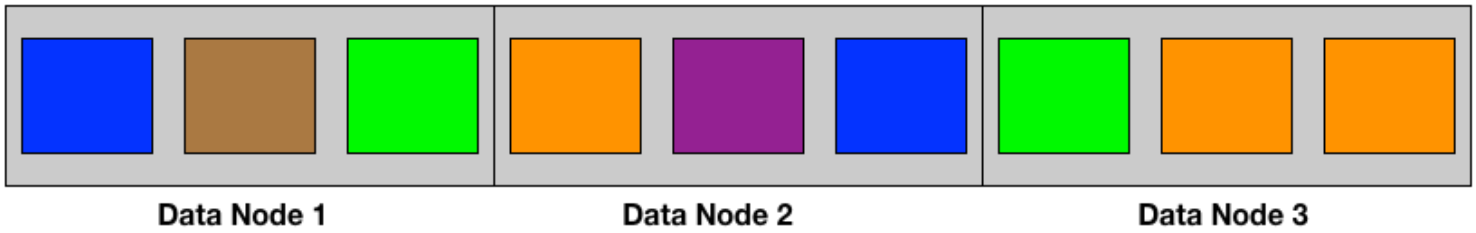

Can we *approximate* bootstrapping without bringing all the data to one place?

Poisson Resampling

Approximation

What if we independently sample each data point using a $Binomial(n, \frac{1}{n})$? All sampling can be done in parallel. For large n, this is *close-enough* to multinomial sampling that it doesn't matter in practice. But, we still need to know n in advance!

Poisson Distribution

$$\lim_{n \to \infty} Binomial(n, \frac{1}{n}) = Poisson(1)$$

$Poisson(\lambda) = \frac{\lambda^k e^{-\lambda}}{k!}$ doesn't need to know n!

Poisson Resampling in Action

Poisson Resampling in Map-Reduce

- Independently sample in your map task using $Poisson(1)$
- Emit those k samples
- Reducers get independent datasets, run their aggregations, and return back results
  - Aggregations could be statistics, or even entire models!


thank you
danke
謝謝
ngiyabonga
teşekkür ederim
gracias
dziękuję
bedankt
obrigado
go raibh maith agat
merci
sukriya
kop khun krap
terima kasih
grazie
arigatõ
dakujem
takk
tack
dank je
спасибо
рахмат
Баярлалаа
faafetai lava
kiitos
dankie
dhanyavad
nanni
nandri
hvala
mauruuru
köszönöm
enkosi
gracie
bayarlalaa
sobodi
dekuji
sagolun
mèsi
didi madloba
najis tuke
kam sah hamnida
তোমাকে ধন্যবাদ
rahmat
chnorakaloutioun
gratias ago
gracies
akun
dankon
aciü
vinaka
спасибі
blagodaram
spas
welalin
kia ora
barka
mersi
misaotra
matondo
paldies
grazzi
mahalo
tapadh leat
хвала
asante
manana
obrigada
tenkí
murakoze
chokrane
djiere dieuf
tau
mochchakkeram
дякую
mamnun
sulpáy
taiku
ありがとう
tanemirt
rahmet
najis tuke
diolch
dhanyavadagalu
shukriya
merce
мерси
ευχαριστώ
감사합니다
xiexie